

# Application design

## Mullins chapter 5

Bjarte Kileng

HVL

January 16, 2018

- 1 Database Application Development and SQL
- 2 Transactions
- 3 Locking
- 4 Isolation Level
- 5 Lock escalation
- 6 Batch processing

# Application performance

## Application performance

Poorly constructed and formulated application code accounts for the majority of relational database performance problems.

## SQL execution time

Different, but logically equivalent SQL can have different execution time.

# SQL – *Structured Query Language*

- ▶ SQL is declarative – We write what we want, not how to retrieve it.
- ▶ SQL versions:
  - (1986~87) ANSI/ISO-standard
  - (1989) SQL-89 = “SQL1”
  - (1992) SQL-92 = “SQL2”
  - (1999) SQL:1999 = “SQL3”
  - (2003) SQL:2003
  - (2006) SQL:2006
  - (2008) SQL:2008
  - (2011) SQL:2011
- ▶ In general, DBMSs follow SQL standards.
- ▶ But, all commercial implementations are different.
- ▶ MySQL – See [MySQL Standards Compliance](#)

```
SELECT /*! STRAIGHT_JOIN */ col1 FROM table1 , table2 WHERE ...
```

# Access Path

- ▶ The DBMS analyzes each SQL and find data-navigational instructions – *access paths*.
- ▶ An *access path* is the “path” the system uses to retrieve and store data in a table.
- ▶ There can be more than one access path available for the same data.
- ▶ Choice of access path has no effect on the semantics (or meaning) of a statement.
- ▶ Choice of access path can have a major effect on the execution time of a statement.

# Set-at-a-Time processing

- ▶ SQL queries and data manipulations are performed on sets of data:
  - A single query return columns and rows – a table.
  - One data manipulation can modify multiple rows.
- ▶ Most programming languages act on one data record at a time.

## Impedance mismatch

Programs expect data to be returned one row at a time, but SQL returns a data set at a time.

# Cursor

- ▶ Most DBMSs provide a *cursor*.
- ▶ Takes input from SQL request.
- ▶ Gives a mechanism to fetch individual rows of a result set.

# Embedding SQL in a program

- ▶ Programming languages need an interface for issuing SQL.
- ▶ Some programming languages support *Embedded SQL*, but most need a API (*connector*).
- ▶ Database abstraction layer:
  - Connector that can talk to many different DBMSs.
  - Provides routines to allocate resources, execute SQL, control DBMS connections, handle transactions, etc.
- ▶ Object-relational mapping (*ORM*, *O/RM*, *O/R mapping*):
  - Connector that map database data with objects in the program.
  - E.g. Java Persistence API (*JPA*), Hibernate.



# Some database abstraction layers

- ▶ ODBC – Open Database Connectivity, originally from Microsoft.
  - Available for many programming languages and DBMSs.
- ▶ JDBC – Java Database Connectivity.
- ▶ MDB2 – PHP.
- ▶ ADO – ActiveX Data Objects.

# Database abstraction layers and drivers

- ▶ A database abstraction layer needs drivers to connect to a DBMS.
- ▶ Optimised for a specific DBMS.
- ▶ MDB2 – Drivers for Firebird, Frontbase, Interbase, MS SQL, MySQL, Oracle, PostgreSQL, Querysim, SQLite.

# Embedded SQL

- ▶ SQL statements are written inline with the program.
- ▶ Pro\*C (Oracle, Sybase) and ECPG (PostgreSQL) example:

```
{  
    int a;  
    /* ... */  
    EXEC SQL SELECT salary INTO :a  
        FROM Employee  
        WHERE SSN=876543210;  
    /* ... */  
    printf("The salary is %d\n", a);  
    /* ... */  
}
```

- ▶ jOOQ embeds a SQL-like language into Java.

# Connectors and DBMSs

- ▶ DBMSs vendors usually provide connectors to their DBMSs:
  - MySQL, see [MySQL::MySQL Connectors](#).
- ▶ Connectors also for specific programming languages:
  - *libdbi* – Database abstraction layer for C.
  - PHP – Include connectors to most DBMSs.
  - Also database abstraction layers for Perl, Python, Java (JDBC).
  - Popular APIs are often copied to other programming languages:
    - *Creole* for PHP – Based on JDBC.
    - *ADOdb* for PHP – ADO for PHP.
- ▶ Client to server communication for Sybase and MS SQL requires FreeTDS on Linux.

# Object orientation and SQL

- ▶ Database – Normalisation and shared data.
- ▶ OO – Encapsulation of data.
- ▶ OO and RDBMS are not inherently compatible.

# Objects and relational database

- ▶ Serialization of object – Storing data using a flat file representation.
  - Also named marshalling.
- ▶ XML representation of data – Supported by many DBMSs.
- ▶ Object-relational mapping (ORM) – JPA, Hibernate.
- ▶ Object-relational database (ORD) – Objects, classes and inheritance are directly supported in schemas and in the query language.
  - Also named object-relational database management system (ORDBMS)

# Types of SQL

- ▶ Planned or unplanned.
- ▶ Embedded or standalone.
- ▶ Dynamic or static.

# Types of SQL – Planned or unplanned

**Planned:** SQL designed and tested for accuracy en efficiency before run in production.

- ▶ Typically embedded into application.

**Unplanned:** Created “on the fly” by end users during course of business.

- ▶ Also called *ad hoc SQL*.
- ▶ Significant source of inefficiency.



# Types of SQL – Embedded or stand-alone

**Embedded:** Embedded within an application program.

**Stand-alone:** SQL run by itself within a query or reporting tool.

# Types of SQL – Dynamic or static

**Dynamic:** Optimized at run time.

- ▶ The SQL depend on input data and the SQL is created at run time.
- ▶ Prepared statements with input parameters are dynamic. Optimization is done only once, but at run time.

**Static:** Optimized prior to execution.

- ▶ The full SQL statement is known at compilation.
- ▶ Can not change without reprogramming and recompilation.
- ▶ Performance of static SQL is generally better than dynamic SQL.

# Questions

- ▶ Can we have static SQL with JDBC?

No. There is no compile-time checking of SQL syntax in Java. Solutions exist though for combining static SQL with Java, e.g. [pureQuery](#) for DB2.

- ▶ Do you see any possibilities for combining static SQL with JDBC?

We can use stored procedures. SQL in stored procedures may be optimized when stored in the DBMS before being run.

# SQL Coding for Performance

## SQL coding for performance

Let the SQL, rather than the application program do the work.

- ▶ Filter out data using the *WHERE* clause of SQL rather than copying data to the application for filtering by the application logic.
- ▶ The more filtering that is done by the DBMS, the less data has to be moved to the application.
- ▶ More later (chapter 12).

# XML, XQuery and SQL/XML

- ▶ Major DBMSs support storing and managing XML data.
- ▶ XQuery – Querying and transforming XML data.
- ▶ XPath – Quering XML data.
- ▶ XSLT – Transforming XML data.
- ▶ SQL/XML – Extension to SQL standard.
  - XML data type.
  - Functions and routines for accessing and manipulating XML data in SQL databases.

# Transactions

- ▶ An atomic unit of work for the DBMS:
  - Several SQL sentences can be put together and run as a unit.
  - No other concurrent connection will make the data inconsistent.
- ▶ When all the instructions have been accomplished, *COMMIT*:
  - All steps since last commit are externalised to the database.
- ▶ *ROLLBACK* – Will move transaction back to state as before the transaction was started.
- ▶ Can save an intermediate state of the transaction using *SAVEPOINT*.
- ▶ Can rollback transaction to a savepoint.
- ▶ MySQL and MariaDB – The XtraDB, InnoDB, NDB, solidDB, IBMDB2I, PrimeBase XT and Falcon engines support transactions.

# ACID (Atomicity, consistency, isolation, durability)

**Atomicity:** All or none of the instructions of the transaction happen.  
The instructions is done as a single unit.

**Consistency:** The consistency of the state of data is preserved:

- ▶ Database go from one valid state to another.

**Isolation:** No transaction sees the other transactions.

- ▶ For a transaction it looks like it have the database alone.
- ▶ Many transactions can run at the same time.
- ▶ Requires a locking mechanism.

**Durability:** A committed transaction should withstand any kind of system failure.

# More on ACID


- ▶ ACID transactions:
  - Will lock shared resources.
  - Should be short in duration.
  - Should never wait for user input when processing a transaction.
- ▶ ACID transactions and locking:
  - Due to the “isolation” requirement.
  - “Isolation” requirement is the most often relaxed ACID requirement.
- ▶ MySQL – Full ACID with the engines XtraDB, InnoDB, solidDB, IBMDB2I, PrimeBase XT, Falcon (MySQL 6).



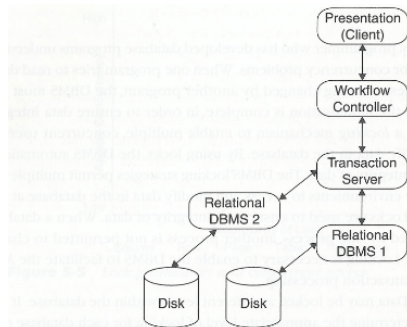
# Transaction Processing Systems

## Transaction server

- ▶ Transaction control on level above DBMSs:
  - Still also transaction control in DBMSs.
- ▶ Why:
  - Environment with multiple, heterogeneous databases.



**Figure 5-3** Using a database server



**Figure 5-4** Using a transaction server

# Transaction Processing Systems

## Application server

- ▶ Include functionality to build, manage and distribute database applications.
- ▶ Usually includes features of a transaction server.
- ▶ Examples:
  - Java – *WebSphere Application Server* (IBM), *Apache Geronimo*, *Glassfish Application Server*, *JBoss* (Red Hat), *WebLogic Server* (Oracle).
  - PHP – Zend.
  - .Net – Base4, Spring Framework (Spring Framework also for Java).
  - Python – Zope.


# Locking

## Locking

When a user updates data and several concurrent users access the same data, locking is required in order to ensure data integrity.

- ▶ Can lock at different levels:
  - Lock granularity.
  - Column, row, page (block), table, tablespace or database locks.
- ▶ The finer the lock granularity:
  - The more concurrent access will be allowed.
  - The more resources are used by the DBMS to handle the locks.
- ▶ Observe that also gaps might be locked:
  - In XtraDB and InnoDB, if a row is locked, also the gap in indexes till the next row will normally be locked.

# Locking granularity



**Figure 5-5** *Lock granularity and concurrent access*

# MySQL and MariaDB

- ▶ All engines have the “LOCK TABLE”.
- ▶ Of the “common” engines, only XtraDB, InnoDB and NDB have row-level locking.
  - The engines solidDB, IBMDB2I, PrimeBase XT and Falcon are not in my “common” group.
- ▶ For InnoDB, see the MySQL manual “[The InnoDB Transaction Model and Locking](#)”.

# Types of locks

- ▶ Shared lock – Concurrent reads. No updates are allowed.
- ▶ Exclusive lock – Data is modified. No other client is granted access.
- ▶ Update lock:
  - Data must be read before it is changed or deleted.
  - Indicates that data may be changed in the future.
  - Legal – One update + multiple shared locks
  - Illegal – Update + (update or exclusive).
  - Concurrent users can read, but not change data.
  - When data is changed, the update lock is changed to an exclusive lock.
- ▶ MySQL/MariaDB – A shared lock can be updated to an exclusive lock if there are no conflicting locks from concurrent processes.

# Intent locks

- ▶ Lock on table (or part of table).
- ▶ Needed when locks at different levels can cause conflicting locks:
  - One user locks some rows in a table. An intention lock is then put on the table as whole.
  - Another client tries to lock the whole table. This is refused due to the table intent lock.
- ▶ MySQL, see the manual “[InnoDB Lock Modes](#)”.

# Lock timeouts

- ▶ Data is locked for update by one client, then other clients must wait for the lock to be released.
- ▶ Clients will not wait forever:
  - The lock timeout value determines how long the client will wait.
  - When timeout occurs, transaction fails.
  - Application should then retry the transaction.
- ▶ MySQL – Parameter “[innodb\\_lock\\_wait\\_timeout](#)”.



# Deadlocks

- ▶ Known from DAT103.
- ▶ Circular wait.
- ▶ DBMS will chose one process to abort and rollback.
- ▶ If deadlock, application should retry the transaction.

# How to avoid deadlocks

- ▶ Change the lock granularity.
- ▶ Access tables and rows in a given order (avoiding circular wait).
- ▶ Use a different isolation level (more on this next).

# Deadlocks in MariaDB and MySQL

- ▶ XtraDB and InnoDB has deadlock detection. See section MySQL “[Deadlock Detection and Rollback](#)”.
- ▶ A deadlock involving InnoDB or XtraDB tables together with tables from other engines may not be detected.
- ▶ Cause of deadlock can be found using:

```
SHOW ENGINE INNODB STATUS
```

- ▶ The above command requires the *SUPER* privilege.

# Minimise Locking Problems

- ▶ Standardise the sequence of updates:
  - Will avoid deadlocks.
  - The deadlock condition *Circular wait* will then not occur.
- ▶ Do all data modification requests at the end of the transaction:
  - Can reduce the lock duration.

# Isolation Levels

- ▶ ACID only with isolation level *SERIALIZABLE*.
  - The other levels break the isolation requirement.
- ▶ The levels are, from lowest to highest isolation:
  - 1 *READ UNCOMMITTED*,
  - 2 *READ COMMITTED*,
  - 3 *REPEATABLE READ*,
  - 4 *SERIALIZABLE*.

## Extract from the Oracle documentation

From *Introduction to Data Concurrency and Consistency*.

In short, real-world considerations usually require a compromise between perfect transaction isolation and performance.

# Setting the isolation level

```
SET [SESSION | GLOBAL] TRANSACTION ISOLATION LEVEL  
{READ UNCOMMITTED | READ COMMITTED | REPEATABLE READ | SERIALIZABLE}
```

# MySQL engines and isolation levels

- ▶ XtraDB, InnoDB and Falcon (MySQL 6) have all levels.
- ▶ Default isolation level in InnoDB is *REPEATABLE READ*.
  - Other DBMSs often has *READ COMMITTED* as the default level.
- ▶ NDB support only the *READ COMMITTED*.
- ▶ The other frequently used engines lack support for transactions.
- ▶ solidDB, IBMDB2I, PrimeBase XT and Falcon are full ACID, but I do not know the details of these engines.

# READ UNCOMMITTED

- ▶ Also named *dirty read*.
- ▶ Has the highest level of availability and concurrency.
- ▶ Worst degree of data integrity.
- ▶ Will read data changed by other transactions still not committed.
- ▶ Only to be used when integrity problems can be tolerated:
  - Finding averages, making statistics, data warehouse.



# READ COMMITTED

- ▶ Also called *cursor stability*.
- ▶ Default level for many DBMSs.
- ▶ A transaction can read data committed by another transaction.

# REPEATABLE READ

- ▶ Data read in the transaction will not change during the transaction.
- ▶ Can be implemented using shared locks.
- ▶ No range locks.
- ▶ Phantom reads are possible.

# Phantom reads

- ▶ Two identical reads in a transaction can return different rows.
- ▶ Transaction  $T_1$  read a range of rows and acquire shared locks on the rows, but no range locks.
- ▶ Transaction  $T_2$  insert a new row inside the range already read by  $T_1$ .
- ▶ Transaction  $T_1$  issues the first read once more.
- ▶ Now, the resultset has changed even though the two reads by  $T_1$  were identical. The new row inserted by  $T_2$  will now appear in the new resultset.

# REPEATABLE READ with MySQL and MariaDB

- ▶ The default isolation level.
- ▶ MySQL uses a kind of snapshot isolation.
- ▶ The first read of data creates a snapshot of the read data.
- ▶ Following reads of the same data will read from the snapshot.
- ▶ The MySQL and MariaDB implementation do probably not conform to the ANSI/ISO SQL standard for *REPEATABLE READ*.
- ▶ *REPEATABLE READ* closer to the ANSI/ISO SQL standard is probably obtained through *SERIALIZABLE* if using:

```
innodb_locks_unsafe_for_binlog=1
```

See the MySQL manual for “[innodb\\_locks\\_unsafe\\_for\\_binlog](#)”.

# SERIALIZABLE

- ▶ Highest isolation.
- ▶ Removes the possibilities for phantoms.
- ▶ Require locks on data and data ranges.
- ▶ Will guarantee that one transaction will finish without any other transaction changing the data used by the transaction.

# SERIALIZABLE with InnoDB and XtraDB

- ▶ All SELECTs are transformed to:

```
SELECT ... LOCK IN SHARE MODE
```

- ▶ MariaDB and InnoDB lock implementation will acquire the necessary range locks to avoid phantom reads.
- ▶ See the MySQL manual, “[InnoDB Record, Gap, and Next-Key Locks](#)”.

# SERIALIZABLE in Oracle and PostgreSQL

- ▶ Uses snapshot isolation.
- ▶ Has similarities with the snapshot model of MySQL and MariaDB for *REPEATABLE READ*.
- ▶ The transaction will successfully commit only if the updates do not conflict with updates from other transactions on the same data
- ▶ Not truly serializable.
- ▶ Can give inconsistencies between concurrent transactions, e.g. the *Write Skew Anomaly*.
  - PostgreSQL is protected against the simple *Write Skew Anomaly* ([ref](#)) but with more than two concurrent transactions, problems can still occur ([ref](#)).
  - See also the [Oracle documentation](#) on *Data Concurrency and Consistency*.

# Locks and resources

## Locks and resources

With many locks, the DBMS might spend much resources (storage) to handle the locks.



# Using lock escalation

- ▶ Increase the lock granularity:
  - E.g. locks on rows are replaced with one lock on the whole table.
- ▶ Will save space since fewer locks have to be stored.
- ▶ Will lock more data:
  - Concurrent access will suffer.

# Implementation in DBMSs

- ▶ Lock escalation is used by some DBMSs:
  - DB2, Microsoft SQL Server.
  - System parameters determine how lock escalation is done.
- ▶ MariaDB, MySQL and Oracle – No automatic lock escalation:
  - Can be done by the application.
  - InnoDB – Lock escalation is not needed due to efficient storage of locks, see [“The InnoDB Transaction Model and Locking”](#).

# Batch processing

- ▶ No online interaction.
- ▶ Do *COMMIT* periodically:
  - Will release locks and increase concurrency availability.
  - If batch program fails, can restart from the last *COMMIT*.
- ▶ Must be easy to restart the program from the last *COMMIT*:
  - The program must track its own progress.
- ▶ Schedule the batch jobs to run during off-peak online processing hours.