

U1 – Chapter 2

Bjarte Kileng

HVL

January 9, 2018

Boot process and configuration

Select boot device:

- ▶ Enable disk as boot device.
- ▶ Choose boot device.
- ▶ Specify boot device priorities.

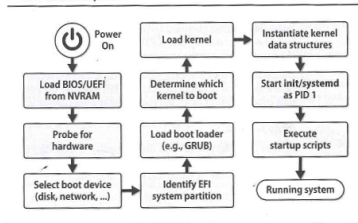
Determine which kernel to boot:

- ▶ Select kernel to load..
- ▶ Specify default kernel.

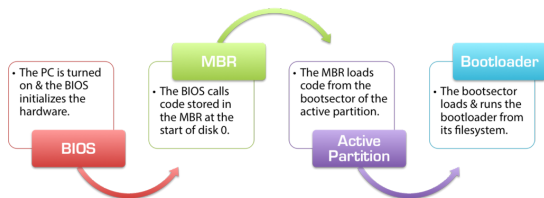
Execute startup scripts:

- ▶ Specify what init scripts to run.
- ▶ Run order or dependencies between init scripts.

Linux & UNIX boot process



- ▶ BIOS and UEFI are firmware that is run when the computer starts.
 - CPU is hardwired to run the firmware.
- ▶ Usually stored in EEPROM on the motherboard.
- ▶ Initializes the computer and runs the next stage of the bootstrapping code (e.g. GRUB).
- ▶ Typically includes a GUI that let us e.g. select the boot disk, or specify a disk priority for boot.
- ▶ UEFI has replaced BIOS on modern PCs.
- ▶ UEFI firmwares usually implement some kind of BIOS compatibility.
 - UEFI can boot from a MBR.



- ▶ Boot disk must start with MBR that contains:
 - First-stage boot loader, *boot block*.
 - Disk partition table.
- ▶ The second-stage boot loader, e.g. GRUB:
 - Includes the necessary file system drivers to load the OS.
 - Loads the OS.
- ▶ The second-stage boot loader is read either from:
 - The active partition, or
 - the dead zone before the first partition (64 disk blocks).
- ▶ All BIOS configurations are done inside the firmware.
 - BIOS has no concept of OS or file systems.

- ▶ UEFI = Unified Extensible Firmware Interface
- ▶ GUID = Globally Unique Identifier
- ▶ GPT = GUID Partition Table
- ▶ ESP = EFI System Partition
- ▶ FAT = File Allocation Table (MS_DOS file system)
- ▶ NVRAM = Non-Volatile RAM

- ▶ UEFI firmware can handle FAT partitions.
- ▶ ESP is a FAT partition that stores the target application, e.g. GRUB.
 - Usually mounted as «/boot/efi».
- ▶ UEFI parameters are stored in NVRAM, accessible by UEFI and OS.
 - Accessible as «/sys/firmware/efi/efivars», or
 - «/sys/firmware/efi/vars».
- ▶ GPT identifies the ESP and the target application.
 - Stored as UEFI parameters.
 - Default target application is «/efi/boot/bootx64.efi».
- ▶ UEFI configurations can be modified in user space using e.g. the command **efibootmgr** or through the «/sys/firmware/» file system.

- ▶ Default boot loader for Linux.
- ▶ Can boot multiple operating systems.
- ▶ GRUB code location:
 - UEFI – ESP
 - BIOS – Usually in the dead zone before the first partition.
- ▶ GRUB can read its configurations from a Linux file system.
 - Usually found as «/boot/grub2/grub.cfg».
- ▶ Modifications in «/boot/grub2/grub.cfg» do not survive updates.
 - Modify instead «/etc/default/grub».
 - Regenerate «/boot/grub2/grub.cfg» with the command:

```
grub2-mkconfig -o /boot/grub2/grub.cfg
```

- Before using *grub2-mkconfig*, «os-prober» must be installed.
- ▶ To specify a default kernel, see [GRUB 2 - Fedora Project Wiki](#).

- ▶ Let us modify the boot entries at boot time.
 - Changes are not saved.
- ▶ See the book or Internet for details.
- ▶ Forgotten the root password? Tell GRUB to start a bash shell:
 - Locate the menu entry.
 - Press the «e» key.
 - Add *init=/bin/bash* to the kernel line.
 - Use «Ctrl-x» to boot.

- ▶ Started by the kernel as process with pid equal to 1.
- ▶ Starts and stops system services and daemons.
- ▶ The systemd process is the top process in the process hierarchy.
 - The Linux kernel can namespace the process hierarchy.
 - Inside a process namespace, pid 1 is the first process of the namespace.

- ▶ The init system of System V is now replaced with systemd.
 - CentOS 7 is the first CentOS to use systemd.
- ▶ Startup sequence:
 - init – the scripts are run in a set sequence.
 - systemd – dependencies determine the startup sequence.
- ▶ Parallelism:
 - init – the scripts are run in a strict sequence, one at a time.
 - systemd – can run the scripts in parallel.
- ▶ systemd is more complicated system.

systemd units and unit files

- ▶ A unit is an entity that is managed by systemd:
 - Can be a service started by a script.
 - Can be a target (more later).
 - And much more.
- ▶ Every systemd unit must have a unit file:
 - Path of executable.
 - Specify how to start and stop.
 - Specify dependencies.
- ▶ With **systemd-run**, units can be started without a unit file.
 - Are named transient units.
- ▶ Unit files are read from:
 - «/etc/systemd/system» – Highest priority.
 - «/usr/lib/systemd/system»
 - «/lib/systemd/system»
- ▶ Unit-file suffix specify unit type:
 - E.g. «.service», «.target», «.timer», «.socket», «.mount», . . .

Find systemd units

- ▶ List all loaded units:

```
systemctl list-units
```

- ▶ List loaded services only:

```
systemctl list-units --type=service
```

- ▶ List all service unit files:

```
systemctl list-unit-files --type=service
```

Enable/disable/mask

- ▶ An enabled service will be started at boot by systemd.

```
systemctl enable NetworkManager-wait-online.service
```

- ▶ A disabled service will not be started at boot by systemd.
 - Can still be started manually.
 - Can also be started due to a depending enabled service.

```
systemctl disable NetworkManager-wait-online.service
```

- ▶ A masked service can not be started.
 - Unit file is linked to «/dev/null».

```
systemctl mask NetworkManager-wait-online.service
```

Start/stop/status

- ▶ A service can be started.

```
systemctl start chronyd.service
```

- ▶ A service can be stopped.

```
systemctl stop chronyd.service
```

- ▶ We can check the status of a service.

```
systemctl status -l chronyd.service
```

or

```
systemctl status -l --no-pager chronyd.service
```

- ▶ Only unit files with an `Install` section can be enabled or disabled with **systemctl**.
- ▶ The **list-unit-files** sub command list service and enablement state.
 - See the man page for all the enablement states.
- ▶ Static services can only be started by hand, or if named as a dependency by an enabled service.

- ▶ Group together unit files.
- ▶ List loaded targets:

```
systemctl list-units --type=target
```

- ▶ List all targets:

```
systemctl list-unit-files --type=target
```

- ▶ Some targets have special meaning.
 - Run automatically by system at specific events.

Some special targets

`ctrl-alt-del.target`: Run when Control+Alt+Del is pressed.

`default.target`: Default unit to start at boot.

`emergency.target`: Run if boot fails, e.g. due to a failing local disk mount.

▶ Play with this mode **before** you need it.

`graphical.target`: Sets up the graphical login screen.

`multi-user.target`: Sets up a non-gui multi-user system.

`reboot.target`: Shutdown and reboot the system.

`shutdown.target`: Shutdown the system.

`rescue.target`: Pulls in the base system for administrative purpose.

Using targets

- ▶ The sub command **isolate** changes the current mode:

```
systemctl isolate multi-user.target
```

- Activate the target with its dependencies, and deactivate all others.

- ▶ Show the default target:

```
systemctl get-default
```

- ▶ Set the default target:

```
systemctl set-default multi-user.target
```

- ▶ Can also specify a target on the kernel boot line
 - Emergency mode: `systemd.unit=emergency.target`
 - Rescue mode: `systemd.unit=rescue.target`.

- ▶ Dependencies are specified in the unit files.
- ▶ Keywords are e.g. *Wants*, *WantedBy*, *Requires*, *RequiredBy*, *Conflicts*.
 - See the book and manual pages for more keywords and their meaning.
 - Does not imply any sequence for processing.
- ▶ Keywords to serialize processing are e.g. *After*, *Before*.
 - If no serializing keyword, systemd will try to run processes in parallel.
- ▶ Systemd assumes a default set of dependencies for units.
 - Turn off assumption with `DefaultDependencies=false`

- ▶ New unit files can be created in an editor.
- ▶ The systemctl sub command **edit** lets us modify an existing unit-file.

```
systemctl edit httpd.service
```

- Creates an override in «/etc/systemd/system/httpd.service.d/».
- ▶ Do not modify existing unit-files.
 - Use an override, e.g.
«/etc/systemd/system/httpd.service.d/override.conf».
- ▶ Using unit-file overrides:
 - The original and the override will be merged by systemd at use.
 - If collisions, the override has higher priority.
- ▶ Install section can not be modified by an override.

Some words on logging

- ▶ Systemd has its own logging system, managed by the *journald* daemon.
- ▶ Systemd log messages are stored in the «/run» directory.
 - Typically, the rsyslog daemon will process also the systemd logs.
- ▶ Can be accessed with the command **journalctl**.
 - Show log of the Bluetooth unit:

```
journalctl -u bluetooth.service
```

- ▶ System logging will be covered later.